

# KickStarterProject

(Eliza) Zijin Huang, Tian Xia

4/17/2019

## 1.

**\*\*Data Acquisition and Integration**

```
data_source_1 <- read.csv("Kickstarter001.csv", header = TRUE, sep = ",")
data_source_2 <- read.csv("Kickstarter002.csv", header = TRUE, sep = ",")

raw_data <- rbind(data_source_1, data_source_2)
```

## 2.

**\*\*Data Cleaning** There are 3784 data totally and there are 3680 projects are completed.

```
live_data <- raw_data %>% filter(raw_data$state == "live")
```

To clean the “catagory” column

```
raw_data$category <- raw_data$category %>%
  str_extract("slug\\":\\\".+\\\",\\\"") %>%
  str_replace_all("\\\",\\\"", "") %>%
  str_replace_all("slug\\":\\\"", "") %>%
  str_replace("/.+", "")
```

To clean the “location” column

```
raw_data$location <- raw_data$location %>%
  str_extract("name\\":\\\".+\\\",\\\"") %>%
  str_replace("\\\",\\\".+\\\", \"") %>%
  str_replace_all("name\\":\\\"", \"")
```

To get rid of creator,photo, slug, urls column

```
raw_data$creator <- NULL
raw_data$photo <- NULL
raw_data$slug <- NULL
raw_data$urls <- NULL
```

To add a preparation\_duration column

```
raw_data$preparation_duration <- raw_data$launched_at - raw_data$created_at
raw_data$preparation_duration_r <- seconds_to_period(raw_data$preparation_duration)
```

To add a launch\_duration column

```
raw_data$launch_duration <- raw_data$deadline - raw_data$launched_at
raw_data$launch_duration_r <- seconds_to_period(raw_data$launch_duration)
raw_data$launch_duration_r <- day(raw_data$launch_duration_r)
```

To convert epoch seconds to readable time

```
raw_data$created_at_readable <- anytime(raw_data$created_at)
raw_data$deadline_readable <- anytime(raw_data$deadline)
raw_data$launched_at_readable <- anytime(raw_data$launched_at)
raw_data$preparation_duration_r <- NULL
```

Transfer raw data into a new variable

```
clean_data <- raw_data
write.csv(clean_data, "./data/data.csv")
head(clean_data)
```

```
##   backers_count
## 1             30
## 2              0
## 3            102
## 4             22
## 5              2
## 6              7
##
## 1 Experience tea and coffee as it should be in our handmade, fine bone china mugs. Made exclusively
## 2   Playing Roles Outside of Basic Education (P.R.O.B.E)\nThe magazine that highlights extracurricu
## 3                                                     A pilot for
## 4 A film about suicide. The struggles of our modern world taking people to their limit and how commo
## 5               Fusing the technical qualities and accuracy of photography with a digital process to
## 6               A digital, interactive magazine and online community fo
##   category converted_pledged_amount country created_at currency
## 1   crafts             1547         GB 1515610761         GBP
## 2 publishing              0         US 1426362805         USD
## 3 film & video          8101         US 1525106061         USD
## 4 film & video          1566         GB 1519854040         GBP
## 5      art              11         US 1407346285         USD
## 6 publishing           826         US 1411150798         USD
##   currency_symbol currency_trailing_code current_currency  deadline
## 1      £              false             USD 1521190409
## 2      $              true              USD 1429112946
## 3      $              true              USD 1531713540
## 4      £              false             USD 1522443600
## 5      $              true              USD 1410484909
## 6      $              true              USD 1414008752
##   disable_communication friends  fx_rate goal  id is_backing
## 1             false      1.308394  1000 1361161119
## 2             false      1.000000  5000  746509287
## 3             false      1.000000  6000 1402909261
## 4             false      1.308394   400  311541751
## 5             false      1.000000 11000  466957735
## 6             false      1.000000  2000 1471254290
```

```

##   is_starrable is_starred launched_at      location
## 1      false      1517306009      London
## 2      false      1426520946      Columbus
## 3      false      1529070876 St. Petersburg
## 4      false      1519937886      Dorset
## 5      false      1407892909      Ypsilanti
## 6      false      1411416752 San Francisco
##                                     name permissions pledged
## 1 Fine Bone China Ceramic Mugs, Made in England      1111.0
## 2               P.R.O.B.E. Magazine                0.0
## 3               'Merican Wasteland Pilot             8101.0
## 4               Cliff - Feature Film                 1116.5
## 5               Photo to Artwork                    11.0
## 6 Lilah Magazine 1st issue launching Dec 2014        826.0
##
## 1
## 2
## 3
## 4 {"id":3322408,"project_id":3322408,"state":"active","state_changed_at":1523464237,"name":"CLIFF - 1
## 5
## 6
##                                     source_url
## 1               https://www.kickstarter.com/discover/categories/crafts
## 2 https://www.kickstarter.com/discover/categories/publishing/periodicals
## 3 https://www.kickstarter.com/discover/categories/film%20&%20video
## 4 https://www.kickstarter.com/discover/categories/film%20&%20video
## 5 https://www.kickstarter.com/discover/categories/art/digital%20art
## 6 https://www.kickstarter.com/discover/categories/publishing/periodicals
## spotlight staff_pick      state state_changed_at static_usd_rate
## 1      true      false successful      1521190409      1.413819
## 2      false      false  failed      1429112947      1.000000
## 3      true      false successful      1531713540      1.000000
## 4      true      false successful      1522443600      1.390235
## 5      false      false  failed      1410484909      1.000000
## 6      false      false  failed      1414008752      1.000000
##   usd_pledged      usd_type preparation_duration launch_duration
## 1    1570.753 international      1695248      3884400
## 2      0.000 international      158141      2592000
## 3    8101.000 international      3964815      2642664
## 4    1552.197 international      83846      2505714
## 5      11.000      domestic      546624      2592000
## 6    826.000 international      265954      2592000
## launch_duration_r created_at_readable  deadline_readable
## 1               44 2018-01-10 13:59:21 2018-03-16 04:53:29
## 2               30 2015-03-14 15:53:25 2015-04-15 11:49:06
## 3               30 2018-04-30 12:34:21 2018-07-15 23:59:00
## 4               29 2018-02-28 16:40:40 2018-03-30 17:00:00
## 5               30 2014-08-06 13:31:25 2014-09-11 21:21:49
## 6               30 2014-09-19 14:19:58 2014-10-22 16:12:32
## launched_at_readable
## 1 2018-01-30 04:53:29
## 2 2015-03-16 11:49:06
## 3 2018-06-15 09:54:36
## 4 2018-03-01 15:58:06

```

```
## 5 2014-08-12 21:21:49
## 6 2014-09-22 16:12:32
```

### 3.

**\*\*Data exploration and Visualization**


3.1 Summarise the number of projects for each status

```
status_projects <- clean_data %>%
  group_by(clean_data$state) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
head(status_projects)
```

```
## # A tibble: 5 x 2
##   `clean_data$state` count
##   <fct>             <int>
## 1 successful         4172
## 2 failed             2801
## 3 canceled           323
## 4 live              253
## 5 suspended          25
```

Plot the number of projects for each status

```
ggplot(status_projects, aes(reorder(status_projects$`clean_data$state`, -count), count, fill=count)) + geom_bar() +
  ggtitle("Projects by Status") + xlab("Project Status") + ylab("Frequency") +
  geom_text(aes(label=count), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```




3.2 Summarise the number of projects for each category

```
catagory_projects <- clean_data %>%
  group_by(clean_data$category) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
head(catagory_projects)
```

```
## # A tibble: 6 x 2
##   `clean_data$category` count
##   <chr>                <int>
## 1 film & video          1102
## 2 music                 858
## 3 technology            795
## 4 art                   787
## 5 publishing            707
## 6 fashion               495
```

Plot the popularity of each category, which is determined by the number of projects

```
ggplot(catagory_projects, aes(reorder(catagory_projects$`clean_data$category`, -count), count, fill=count)) +
  ggtitle("Projects by Category") + xlab("Project Category") + ylab("Frequency") +
  geom_text(aes(label=count), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```




3.3 What types of projects are being funded?

```
pledged_category <- clean_data %>%
  group_by(clean_data$category) %>%
  summarise(total = sum(usd_pledged)) %>%
  arrange(desc(total))
head(pledged_category)
```

```
## # A tibble: 6 x 2
##   `clean_data$category`      total
##   <chr>                  <dbl>
## 1 technology            31373736.
## 2 film & video          10398557.
## 3 design                8108525.
## 4 fashion               6797765.
## 5 games                 6465490.
## 6 music                 4988363.
```

Plot the amount pledged by each category

```
ggplot(pledged_category, aes(reorder(pledged_category$`clean_data$category`, -total), total/1000000, fill=total)) +
  ggtitle("Total Amount Pledged by Category") + xlab("Project Category") +
  ylab("Amount Pledged (USD millions)") +
  geom_text(aes(label=paste0("$", round(total/1000000,1))), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```



3.4 How much is pledged per backer for each category?


```
pledged_avg_category <- clean_data %>%
  group_by(clean_data$category) %>%
  summarise(pledged = sum(usd_pledged), backers=sum(backers_count)) %>%
  mutate(avg = pledged/backers) %>%
  arrange(desc(avg))
head(pledged_avg_category)
```

```
## # A tibble: 6 x 4
##   `clean_data$category` pledged backers  avg
##   <chr>                <dbl>   <int> <dbl>
## 1 technology          31373736.  224188 140.
## 2 design              8108525.   73647 110.
## 3 fashion             6797765.   63802 107.
## 4 film & video       10398557.  104058 99.9
## 5 food               2745202.   29880 91.9
## 6 photography        1620027.   18585 87.2
```

Plot the amount pledged per backer for each category

```
ggplot(pledged_avg_category, aes(reorder(clean_data$category`, -avg), avg, fill=a
  ggtitle("Average Amount Pledged per Backer") + xlab("Project Category") +
  ylab("Amount Pledged (USD)") +
```

```
geom_text(aes(label=paste0("$", round(avg,2))), vjust=-0.5) + theme_economist() +
theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
axis.text.x=element_text(size=12, angle=90), legend.position="null") +
scale_fill_gradient(low="skyblue1", high="royalblue4")
```



3.5 Get the 10 highest goal successful projects

```
top_ten_success <- clean_data[clean_data$state == "successful",] %>%
  select("category", "goal", "state") %>%
  arrange(desc(goal))
head(top_ten_success)
```

```
##      category    goal    state
## 1  technology 1500000 successful
## 2  technology  800000 successful
## 3  technology  800000 successful
## 4 photography  700000 successful
## 5  technology  500000 successful
## 6      design  500000 successful
```

3.6 Get the average project goal

```
goal_avg <- clean_data %>%
  group_by(category) %>%
  summarise(goals = sum(goal), projects = n()) %>%
```




```
mutate(avg = goals/projects) %>%
  arrange(desc(avg))
head(goal_avg)
```

```
## # A tibble: 6 x 4
##   category      goals projects    avg
##   <chr>         <dbl>    <int>  <dbl>
## 1 design      104348985    417 250237.
## 2 film & video 150585489.   1102 136647.
## 3 technology   66183820    795  83250.
## 4 theater      9656489    221  43695.
## 5 journalism   6530680    170  38416.
## 6 food        16506111    458  36040.
```

Plot the average project goal.

```
ggplot(goal_avg, aes(reorder(goal_avg$category, -avg), avg, fill=avg)) + geom_bar(stat="identity") +
  ggtitle("Average Project Goal") + xlab("Project Category") + ylab("Project Goal (USD)") +
  geom_text(aes(label=paste0("$", round(avg,0))), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```



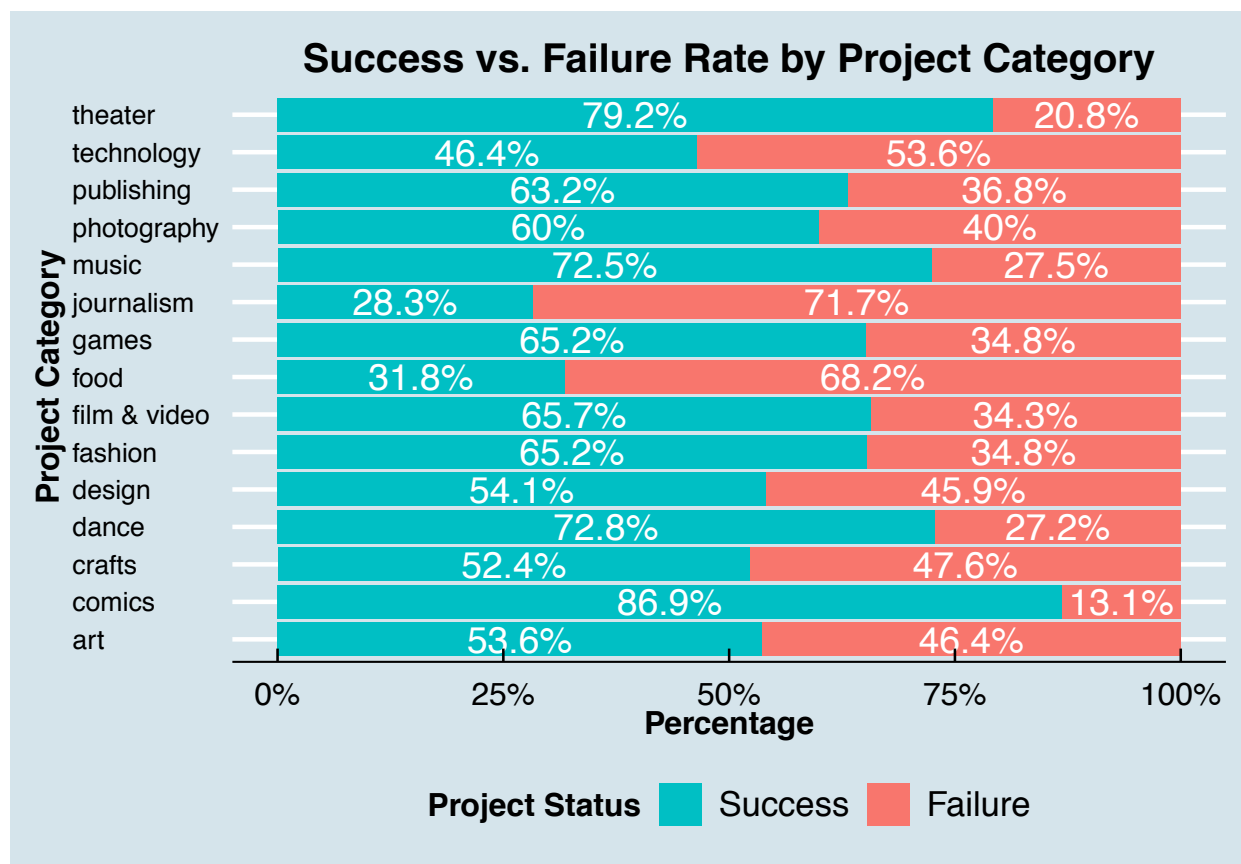
3.7 percentage for projects in each category

```
perc_projects <- clean_data %>%
  filter(state %in% c("successful", "failed")) %>%
  group_by(category, state) %>%
  summarize(count=n()) %>%
  mutate(pct=count/sum(count)) %>%
  arrange(desc(state), pct)
head(perc_projects)
```

```
## # A tibble: 6 x 4
## # Groups:   category [6]
##   category    state      count    pct
##   <chr>      <fct>      <int> <dbl>
## 1 journalism successful    43 0.283
## 2 food        successful   132 0.318
## 3 technology  successful   344 0.464
## 4 crafts      successful   199 0.524
## 5 art         successful   390 0.536
## 6 design      successful   198 0.541
```

Plot the percentage for each category

```
ggplot(perc_projects, aes(perc_projects$category, pct, fill=state)) + geom_bar(stat="identity") +
  ggtitle("Success vs. Failure Rate by Project Category") +
  xlab("Project Category") + ylab("Percentage") + scale_y_continuous(labels=scales::percent) +
  scale_fill_discrete(name="Project Status", breaks=c("successful", "failed"),
    labels=c("Success", "Failure")) +
  geom_text(aes(label=paste0(round(pct*100,1),"%")), position=position_stack(vjust=0.5),
    colour="white", size=5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
    axis.text.x=element_text(size=12), legend.position="bottom",
    legend.title=element_text(size=12, face="bold")) + coord_flip()
```

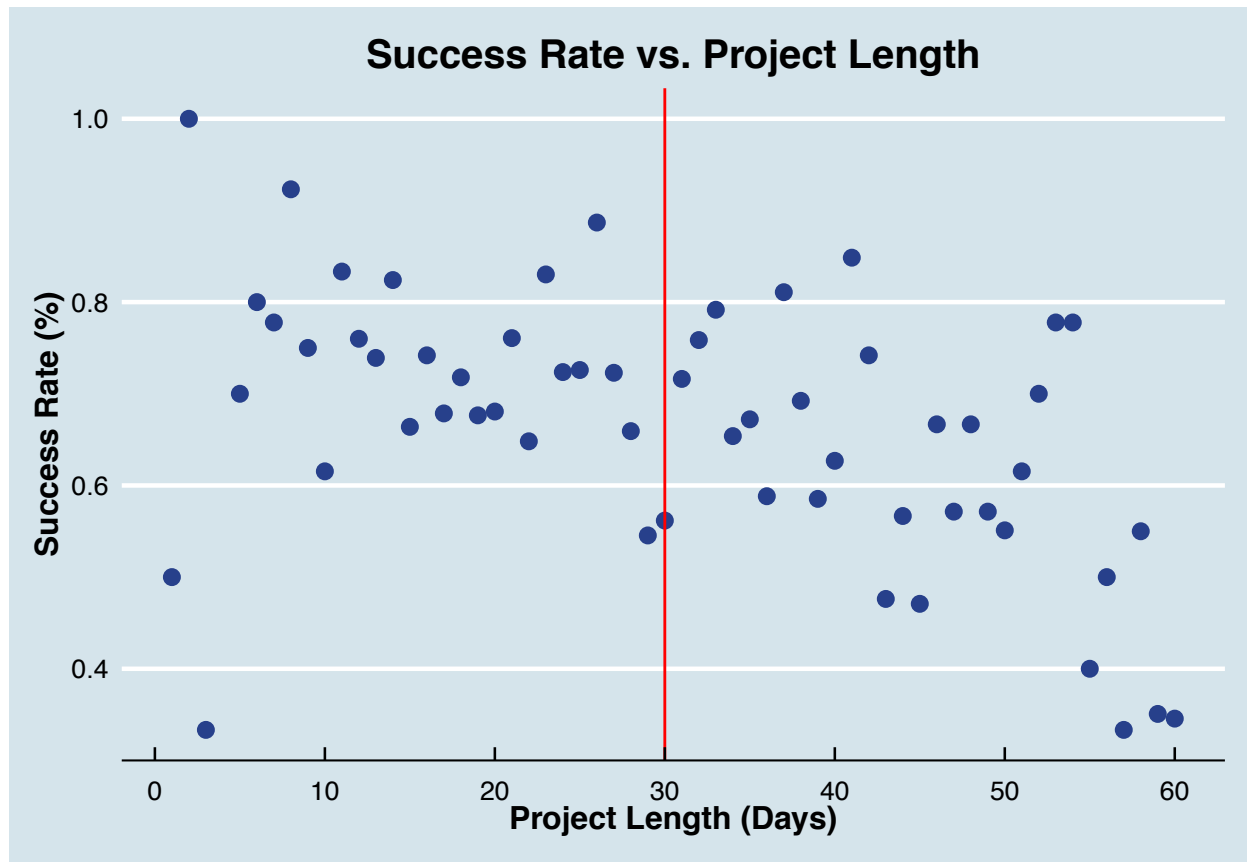


3.8 Does project length affect success rate?

```
perc_length <- clean_data %>%
  filter(state %in% c("successful", "failed"), launch_duration_r < 61) %>%
  group_by(launch_duration_r, state) %>%
  summarize(count=n()) %>%
  mutate(pct=count/sum(count))
head(perc_length)
```

```
## # A tibble: 6 x 4
## # Groups:   launch_duration_r [4]
##   launch_duration_r state      count  pct
##             <dbl> <fct>      <int> <dbl>
## 1             1 failed         1 0.5
## 2             1 successful       1 0.5
## 3             2 successful       1 1
## 4             3 failed         6 0.667
## 5             3 successful       3 0.333
## 6             4 failed         2 1
```

```
ggplot(perc_length[perc_length$state=="successful",], aes(launch_duration_r, pct)) +
  geom_point(colour="royalblue4", size=2.5) + ggtitle("Success Rate vs. Project Length") +
  xlab("Project Length (Days)") + ylab("Success Rate (%)") +
  scale_x_continuous(breaks=c(0,10,20,30,40,50,60)) + geom_vline(xintercept=30, colour="red") +
  theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"))
```



3.9 Where it is coming form?


```
countries_freq <- clean_data %>%
  group_by(country) %>%
  summarize(count=n())

countries.match <- joinCountryData2Map(countries_freq, joinCode="ISO2", nameJoinColumn="country")

## 22 codes from your data successfully matched countries in the map
## 0 codes from your data failed to match with a country code in the map
## 220 codes from the map weren't represented in your data

mapCountryData(countries.match, nameColumnToPlot="count",
  mapTitle="Number of Projects by Country", catMethod="logFixedWidth",
  colourPalette="heat")
```

## Number of Projects by Country



# Data Model Construction and Prediction

April 17, 2019

## 1 Predicting Crowdfunding Success

Using kNN, logistic regression model, support vector machine, naive bayes to predict success rate of crowdfunding.

```
In [1]: import re
import pandas as pd
import numpy as np
import seaborn as sns
import datetime
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import train_test_split

from sklearn.metrics import make_scorer, accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix

from sklearn.naive_bayes import GaussianNB

%matplotlib inline
```

### 1.1 Data Cleaning & Processing

```
In [2]: df = pd.read_csv('./data/data.csv')
df.head(1)
```

```
Out [2]:   Unnamed: 0  backers_count  \
0           1                21
```

blurb category \

```

0 2006 was almost 7 years ago... Can you believ...      Rock

  converted_pledged_amount country created_at currency currency_symbol \
0                802      US 1387659690      USD      $

  currency_trailing_code      ...      static_usd_rate  usd_pledged \
0                True      ...                1.0      802.0

  usd_type preparation_duration preparation_duration_r \
0 international      351356      4d 1H 35M 56S

  launch_duration launch_duration_r created_at_readable \
0      3888000      45d 0H 0M 0S 2013-12-21 16:01:30

  deadline_readable launched_at_readable
0 2014-02-08 17:37:26 2013-12-25 17:37:26

[1 rows x 41 columns]

```

```
In [3]: df.shape
```

```
Out[3]: (3779, 41)
```

```
In [4]: df.state.value_counts()
```

```

Out[4]: successful      2224
        failed         1276
        canceled         149
        live           120
        suspended        10
        Name: state, dtype: int64

```

```

In [5]: # drop status rows labeled as live, canceled, suspended.
df = df[~df['state'].isin(['live', 'canceled', 'suspended'])]
df.shape

```

```
Out[5]: (3500, 41)
```

```

In [6]: # drop irrelevant or independent variables
df.drop(['Unnamed: 0', 'blurb', 'created_at', 'currency_symbol', 'currency_trailing_co
        'deadline', 'disable_communication', 'friends', 'id',
        'is_backing', 'is_starred', 'launched_at', 'state_changed_at',
        'name', 'permissions', 'profile', 'source_url', 'staff_pick',
        'preparation_duration_r', 'launch_duration_r',
        'created_at_readable', 'deadline_readable', 'launched_at_readable',
        'location', 'usd_type'], axis = 1, inplace = True)
df.head()

```

```

Out[6]:  backers_count      category converted_pledged_amount country currency \
0                21      Rock                802      US      USD

```

1	97	Mixed Media	2259	US	USD
2	88	Photobooks	29638	US	USD
3	193	Footwear	49158	IT	EUR
4	20	Software	549	US	USD

	fx_rate	goal	is_starrable	pledged	spotlight	state \
0	1.000000	200.0	False	802.0	True	successful
1	1.000000	400.0	False	2259.0	True	successful
2	1.000000	27224.0	False	29638.0	True	successful
3	1.128433	40000.0	False	43180.0	True	successful
4	1.000000	1000.0	False	549.0	False	failed

	static_usd_rate	usd_pledged	preparation_duration	launch_duration
0	1.000000	802.000000	351356	3888000
1	1.000000	2259.000000	413843	1728000
2	1.000000	29638.000000	769946	2595600
3	1.136525	49075.152523	314662	3625358
4	1.000000	549.000000	212500	2592000

```
In [7]: df['state'] = df.state.str.contains('successful').astype(int)
```

```
In [8]: # add column representing continent
```

```
def classifier(row):
    if row.country in ['US', 'CA', 'GT', 'MX', 'PR', 'NI', 'SV', 'PA', 'BO', 'GU']:
        return 'America'
    elif row.country in ['NG', 'GH', 'ZA', 'KE', 'ET', 'CD', 'MA', 'TZ', 'ZM', 'LR', 'I']:
        return 'Africa'
    elif row.country in ['GB', 'NO', 'DE', 'SE', 'BA', 'IS', 'HU', 'IT', 'NL', 'FR', 'U']:
        return 'Europe'
    elif row.country in ['JM', 'HT', 'BS', 'DO', 'LC', 'DO', 'TT']:
        return 'Carribean'
    elif row.country in ['CN', 'TW', 'HK', 'NP', 'ID', 'SG', 'IN', 'JP', 'LB', 'KZ', 'I']:
        return 'Asia'
    elif row.country in ['IL', 'QA', 'AF', 'KZ', 'AE', 'PS', 'SY', 'SA', 'IQ', 'IR', 'TJ',]:
        return 'Arab'
    else:
        return "Oceania"
```

```
df["continent"] = df.apply(classifier, axis=1)
```

```
In [9]: df.head()
```

```
Out[9]:
```

	backers_count	category	converted_pledged_amount	country	currency \
0	21	Rock	802	US	USD
1	97	Mixed Media	2259	US	USD
2	88	Photobooks	29638	US	USD
3	193	Footwear	49158	IT	EUR
4	20	Software	549	US	USD



	fx_rate	goal	is_starrable	pledged	spotlight	state	\
0	1.000000	200.0	False	802.0	True	1	
1	1.000000	400.0	False	2259.0	True	1	
2	1.000000	27224.0	False	29638.0	True	1	
3	1.128433	40000.0	False	43180.0	True	1	
4	1.000000	1000.0	False	549.0	False	0	

	static_usd_rate	usd_pledged	preparation_duration	launch_duration	\
0	1.000000	802.000000		351356	3888000
1	1.000000	2259.000000		413843	1728000
2	1.000000	29638.000000		769946	2595600
3	1.136525	49075.152523		314662	3625358
4	1.000000	549.000000		212500	2592000

	continent
0	America
1	America
2	America
3	Europe
4	America

```
In [10]: from sklearn import preprocessing
def encode_features(df):
    features = ['category', 'country', 'currency', 'is_starrable', 'continent', 'spotlight']
    df_combined = pd.concat([df, df[features]])

    for feature in features:
        le = preprocessing.LabelEncoder()
        le = le.fit(df_combined[feature])
        df[feature] = le.transform(df_combined[feature])
    return df

data = encode_features(df)
data.head()
```

```
Out[10]:
```

	backers_count	category	converted_pledged_amount	country	currency	\
0	21	120	802	20	13	
1	97	83	2259	20	13	
2	88	99	29638	20	13	
3	193	59	49158	12	4	
4	20	126	549	20	13	

	fx_rate	goal	is_starrable	pledged	spotlight	state	\
0	1.000000	200.0	0	802.0	1	1	
1	1.000000	400.0	0	2259.0	1	1	
2	1.000000	27224.0	0	29638.0	1	1	
3	1.128433	40000.0	0	43180.0	1	1	
4	1.000000	1000.0	0	549.0	0	0	

	static_usd_rate	usd_pledged	preparation_duration	launch_duration	\
0	1.000000	802.000000	351356	3888000	
1	1.000000	2259.000000	413843	1728000	
2	1.000000	29638.000000	769946	2595600	
3	1.136525	49075.152523	314662	3625358	
4	1.000000	549.000000	212500	2592000	

	continent
0	0
1	0
2	0
3	2
4	0

```
In [11]: df.continent.value_counts()
```

```
Out[11]: 0    2677
         2     689
         3     98
         1     36
         Name: continent, dtype: int64
```

```
In [12]: X = df.drop(['preparation_duration', 'launch_duration', 'state', 'backers_count', 'sp
y = df['state']
```

```
In [13]: from sklearn.preprocessing import Imputer
         X = Imputer().fit_transform(X)
```

## 1.2 kNN Model

```
In [14]: k_range = range(1,200)
         k_scores = []
         for k in k_range:
             knn = KNeighborsClassifier(n_neighbors=k)
             scores = cross_val_score(knn, X, y, cv=10, scoring = 'accuracy')
             k_scores.append(scores.mean())
         print('Computed k_scores for k value in range 1 to 200.')
```

Computed k\_scores for k value in range 1 to 200.

```
In [15]: scores.mean()
```


```
Out[15]: 0.9605670786816919
```

```
In [16]: scores.max()
```

```
Out[16]: 0.9885386819484241
```

```
In [16]: plt.plot(k_range, k_scores)
plt.xlabel('value of k for knn')
plt.ylabel('cross-validated accuracy')
```

```
Out[16]: Text(0,0.5,'cross-validated accuracy')
```




```
In [17]: MSE = [1 - x for x in k_scores]
optimal_k = k_range[MSE.index(min(MSE))]
print("The optimal number of neighbors is %d" % optimal_k)
```

```
The optimal number of neighbors is 10
```

### 1.3 Logistic Regression Model

```
In [18]: df.state.value_counts().plot(kind = 'box')
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x1a187ffd30>
```



```
In [19]: ss = StandardScaler()
         lr = LogisticRegression()
         lr_pipe = Pipeline([('sscale', ss), ('logreg', lr)])
```

```
In [20]: lr_pipe.fit(X, y)
```

```
Out[20]: Pipeline(memory=None,
                  steps=[('sscale', StandardScaler(copy=True, with_mean=True, with_std=True)), ('logreg', LogisticRegression(intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False))])
```

```
In [21]: lr_pipe.score(X,y)
```

```
Out[21]: 0.838
```

```
In [22]: # divide the dataset into
         # - 70% training data
         # - 30% test data
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

```
In [23]: lr_pipe.fit(X_train, y_train)
```

```
Out[23]: Pipeline(memory=None,
                  steps=[('sscale', StandardScaler(copy=True, with_mean=True, with_std=True)), ('logreg', LogisticRegression(intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False))])
```

```
In [24]: lr_pipe.score(X_test, y_test)
```

```
Out[24]: 0.8133333333333334
```

```
In [25]: y_pred = lr_pipe.predict(X_test)
```

```
In [26]: from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, c
```

```
In [27]: print(f1_score(y_test, y_pred, average="macro"))
print(precision_score(y_test, y_pred, average="macro"))
print(recall_score(y_test, y_pred, average="macro"))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.7979183032207384
```

```
0.802507012622721
```

```
0.794344333478072
```

```
[[282 110]
```

```
 [ 86 572]]
```

	precision	recall	f1-score	support
0	0.77	0.72	0.74	392
1	0.84	0.87	0.85	658
avg / total	0.81	0.81	0.81	1050

## 1.4 Support Vector Machine

```
In [ ]: svcclassifier = SVC(kernel='linear')
svcclassifier.fit(X_train, y_train)
```

```
In [31]: y_pred = svcclassifier.predict(X_test)
```

```
In [32]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[386  0]
```

```
 [  0 664]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	386
1	1.00	1.00	1.00	664
avg / total	1.00	1.00	1.00	1050

## 1.5 Naive Bayes

```
In [28]: gnb = GaussianNB()
         y_pred = gnb.fit(X_train, y_train).predict(X_test)
```

```
In [29]: print(confusion_matrix(y_test,y_pred))
         print(classification_report(y_test,y_pred))
```


```
[[383   9]
 [479 179]]
```

	precision	recall	f1-score	support
0	0.44	0.98	0.61	392
1	0.95	0.27	0.42	658
avg / total	0.76	0.54	0.49	1050

## 1.6 Data Visualization

```
In [33]: plots = sns.pairplot(df, vars = ['backers_count', 'country', 'goal'],
                                     hue="state")
```

```
/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



```
In [34]: df = df.drop(['is_starrable'], axis=1)
corr = df.corr()
corr
```

```
Out [34]:
```

	backers_count	category	converted_pledged_amount	\
backers_count	1.000000	0.074279	0.855609	
category	0.074279	1.000000	0.056564	
converted_pledged_amount	0.855609	0.056564	1.000000	
country	0.031723	-0.007631	0.031882	
currency	0.027896	-0.011621	0.027261	
fx_rate	-0.023226	-0.007235	-0.008231	
goal	0.330618	0.039287	0.198685	
pledged	0.357206	0.034792	0.229725	
spotlight	0.174372	-0.062161	0.140590	
state	0.174372	-0.062161	0.140590	
static_usd_rate	-0.031860	-0.006960	-0.015069	

usd_pledged	0.855154	0.056879	0.999866
preparation_duration	0.051476	-0.007718	0.061744
launch_duration	0.043758	0.028363	0.051582
continent	-0.023403	0.023315	-0.015755

	country	currency	fx_rate	goal	pledged \
backers_count	0.031723	0.027896	-0.023226	0.330618	0.357206
category	-0.007631	-0.011621	-0.007235	0.039287	0.034792
converted_pledged_amount	0.031882	0.027261	-0.008231	0.198685	0.229725
country	1.000000	0.984185	0.005072	-0.015072	-0.011316
currency	0.984185	1.000000	0.001718	-0.020479	-0.016289
fx_rate	0.005072	0.001718	1.000000	-0.111385	-0.107551
goal	-0.015072	-0.020479	-0.111385	1.000000	0.991735
pledged	-0.011316	-0.016289	-0.107551	0.991735	1.000000
spotlight	0.047634	0.048726	0.002059	-0.000156	0.023665
state	0.047634	0.048726	0.002059	-0.000156	0.023665
static_usd_rate	-0.102557	-0.106477	0.963558	-0.102855	-0.099714
usd_pledged	0.031217	0.026642	-0.008874	0.204915	0.235897
preparation_duration	0.026465	0.028097	-0.002326	0.009646	0.010948
launch_duration	-0.016174	-0.018033	-0.025344	0.020797	0.011791
continent	-0.746953	-0.773095	0.195684	0.015929	0.012159

	spotlight	state	static_usd_rate	usd_pledged \
backers_count	0.174372	0.174372	-0.031860	0.855154
category	-0.062161	-0.062161	-0.006960	0.056879
converted_pledged_amount	0.140590	0.140590	-0.015069	0.999866
country	0.047634	0.047634	-0.102557	0.031217
currency	0.048726	0.048726	-0.106477	0.026642
fx_rate	0.002059	0.002059	0.963558	-0.008874
goal	-0.000156	-0.000156	-0.102855	0.204915
pledged	0.023665	0.023665	-0.099714	0.235897
spotlight	1.000000	1.000000	-0.004343	0.140219
state	1.000000	1.000000	-0.004343	0.140219
static_usd_rate	-0.004343	-0.004343	1.000000	-0.015418
usd_pledged	0.140219	0.140219	-0.015418	1.000000
preparation_duration	0.012206	0.012206	-0.008198	0.061616
launch_duration	-0.144859	-0.144859	-0.025125	0.051350
continent	-0.043037	-0.043037	0.285074	-0.015234

	preparation_duration	launch_duration	continent
backers_count	0.051476	0.043758	-0.023403
category	-0.007718	0.028363	0.023315
converted_pledged_amount	0.061744	0.051582	-0.015755
country	0.026465	-0.016174	-0.746953
currency	0.028097	-0.018033	-0.773095
fx_rate	-0.002326	-0.025344	0.195684
goal	0.009646	0.020797	0.015929
pledged	0.010948	0.011791	0.012159



spotlight	0.012206	-0.144859	-0.043037
state	0.012206	-0.144859	-0.043037
static_usd_rate	-0.008198	-0.025125	0.285074
usd_pledged	0.061616	0.051350	-0.015234
preparation_duration	1.000000	0.029690	-0.026049
launch_duration	0.029690	1.000000	-0.006020
continent	-0.026049	-0.006020	1.000000

In [35]: `sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns)`

Out[35]: `<matplotlib.axes._subplots.AxesSubplot at 0x1a1b3e6358>`

