

KickStarter Project

(Eliza) Zijin Huang, Tian Xia

4/17/2019

Part I. Report

Summary and Problem Statement

A description of the problem that is driving the project (comprehensible by an educated lay person).

In this project, we are predicting the success rate of Kickstarter campaigns. Kickstarter is the world's largest funding platform for creative projects, with projects ranging from technology to arts. Kickstarter's rule states that if a project falls short of meeting its minimum funding goal by its deadline, the project will not receive any fund. Thus very important to set a realistic goal. However, there is currently no practical suggestion or guidance, and the analytics dashboard is hard to set up. Therefore, we are interested in predicting the projects' success rate given a project's goal and other independent variables. Predicting the success rate increases the number of funding goals reached. Ultimately, it supports the entrepreneurs, and in the process support art, advocate for social causes, and bring innovation ideas and products to the world.

Data

A rationale for selecting the data collected used to investigate the problem.

The dataset we used is pre-scraped data stored in multiple csv files. It includes 7,574 projects from North America, Europe, Pacific countries, and Japan. Each project in our dataset has either failed, successful, live, suspended or canceled status. Since we are only interested in projects with status of either successful or failed, we filtered out projects whose status is not successful or failed. In addition we also dropped projects that have missing values.

Our selected dataset is complete and good for us to analyze the data and construct data predictive models. We payed attention to some specific parameters that are important for our analysis such as projects' goal, amount pledged, category, duration, region, etc. All those features are included, even though the data might be untidy or not in the best format for analysis. However, some unrelated parameters such as projects' description, urls, photo have nothing to do with success rate. Thus, we drop those columns and leave parameters that are closely related with project's success.

Cleaning, Visualization, and Prediction

Justification for the processes implemented, and the technical choices you made for your project.

For data cleaning, we chose to use tidyverse, lubridate packages with regular expression and some baseR functions. In order to extract useful information and get rid of irrelevant data, we used regular expression to extract category, location and tidy other parameters into standard format. Also we used lubridate package to create duration columns based on given information about launched date and end date.

In the process of data exploration and visualization analysis, we chose to use ggthemes and rworldmap packages for visualizing the data and dplyr packages to select relevant data. Ggthemes can help us make data more visualized in terms of colors of bars, lines, and points. For example, the higher proportion the category takes or the higher the success rate is, the darker the color of that category is. We also used rworldmap package because we want to present and acknowledge where our data is coming from. It is relatively wide-ranged, even though it does not cover Africa and South America.

For prediction, we used the following models: k-nearest neighbour, logistic regression, support vector machine, and naive bayes. K-nn model is commonly used in classification and regression predictive problems, but has a long calculation time. Simple logistic regression model is used because there are multiple independent variables. SVM is used for classification because it can find an optimal boundary between the classes. NB model takes less time to train, and requires less data. Thus, we choose to implement the above four models for our prediction/classification task.

Issues and Solution

A description of any issues you ran into and how you resolved them.

When we clean the dataset, one challenge we had is that some of columns' data are stored in JSON format and they actually compress a lot of relevant data. The approach we took to solve this problem is by using string extract and replace with regular expression to extract useful information and split them into several columns. In addition, the original data does not have the duration in format of number of days, instead we had launched date and end date, both in Unix time. We used R packages to convert them into readable year/month/day format, and computed two duration variables with the three given timestamp.

In the prediction process, we encounter the following problems:

First, the data frame contains a mix of independent variables and dependent variables. To solve this problem, we implemented a correlation plot of all the variables, and dropped those that are in fact dependent variables. For instance, we dropped "spotlight" (meaning if a Kickstarter project is featured on the homepage of the website), because it is highly correlated with the final state (successful or unsuccessful), and is in fact a dependent variable.

Second, some models (SVM) takes has a long calculation time in R. As a result, we choose to implement the models in python using jupyter notebook for shorter time and more suitable visualization tools (such as Matplotlib).

Please refer to Part II and Part III of this report for input code and session outputs.

Insights and Potential Future Work

Insights on what you learned and potential future work.

Throughout this project, We have learned some things that we do not learn in lectures or homeworks. First, it is our first time to play around with such huge, untidy, and real data. We had some challenges for cleaning data but solved them eventually (refer to the issue section for more details). We realized that a fairly confident prediction can be generated based on our chosen variables. Thus, it is possible to predict the success, and potentially make modification to the project setup, in order to improve the success rate of Kickstarter crowdfunding projects.

In future, we are planning to do a multi-class/multinomial classification or prediction, instead of the binary classification process we are using. This would give more precise feedbacks to the project initiators.

We are also considering building interactive visualization tools (possibly chrome extension) to analyze a specific project's parameters, and give real-time feedback to the project initiators. The feedback could focus on variables such as project duration, pledged amount, areas of improvement on project description etc. Some examples of the suggestions are: "Use a longer time period of 100 days." "Reduce the pledged amount to 1000 dollars." "Replace the word 'must' with 'could'." "Set more pledging options with smaller amounts, such as \$5."

Part II. Data Cleaning and Visualization

#1. **Data Acquisition and Integration

```
data_source_1 <- read.csv("../data/Kickstarter001.csv", header = TRUE, sep = ",")
data_source_2 <- read.csv("../data/Kickstarter002.csv", header = TRUE, sep = ",")

raw_data <- rbind(data_source_1, data_source_2)
```

#2. **Data Cleaning There are 3784 data totally and there are 3680 projects are completed.

```
live_data <- raw_data %>% filter(raw_data$state == "live")
```

To clean the “category” column

```
raw_data$category <- raw_data$category %>%
  str_extract("slug\\":\\\".+\\\",\\\"") %>%
  str_replace_all("\\\",\\\"\", \"\") %>%
  str_replace_all("slug\\":\\\"\", \"\") %>%
  str_replace("/.+\", \"\")
```

To clean the “location” column

```
raw_data$location <- raw_data$location %>%
  str_extract("name\\":\\\".+\\\",\\\"") %>%
  str_replace("\\\",\\\".+\", \"\") %>%
  str_replace_all("name\\":\\\"\", \"\")
```

To get rid of creator,photo, slug, urls column

```
raw_data$creator <- NULL
raw_data$photo <- NULL
raw_data$slug <- NULL
raw_data$urls <- NULL
```

To add a preparation_duration column

```
raw_data$preparation_duration <- raw_data$launched_at - raw_data$created_at
raw_data$preparation_duration_r <- seconds_to_period(raw_data$preparation_duration)
```

To add a launch_duration column

```
raw_data$launch_duration <- raw_data$deadline - raw_data$launched_at
raw_data$launch_duration_r <- seconds_to_period(raw_data$launch_duration)
raw_data$launch_duration_r <- day(raw_data$launch_duration_r)
```

To convert epoch seconds to readable time

```
raw_data$created_at_readable <- anytime(raw_data$created_at)
raw_data$deadline_readable <- anytime(raw_data$deadline)
raw_data$launched_at_readable <- anytime(raw_data$launched_at)
raw_data$preparation_duration_r <- NULL
```

Transfer raw data into a new variable

```
clean_data <- raw_data
write.csv(clean_data, "../data/data.csv")
head(clean_data)
```

```
## backers_count
```

```

## 1          30
## 2          0
## 3         102
## 4         22
## 5          2
## 6          7
##
## 1 Experience tea and coffee as it should be in our handmade, fine bone china mugs. Made exclusively :
## 2   Playing Roles Outside of Basic Education (P.R.O.B.E.)\nThe magazine that highlights extracurricu
## 3                                           A pilot for
## 4 A film about suicide. The struggles of our modern world taking people to their limit and how common
## 5           Fusing the technical qualities and accuracy of photography with a digital process to
## 6                               A digital, interactive magazine and online community fo
##      category converted_pledged_amount country created_at currency
## 1      crafts                1547      GB 1515610761      GBP
## 2    publishing                0      US 1426362805      USD
## 3 film & video            8101      US 1525106061      USD
## 4 film & video            1566      GB 1519854040      GBP
## 5         art                11      US 1407346285      USD
## 6    publishing            826      US 1411150798      USD
##      currency_symbol currency_trailing_code current_currency  deadline
## 1          £                false          USD 1521190409
## 2          $                true          USD 1429112946
## 3          $                true          USD 1531713540
## 4          £                false          USD 1522443600
## 5          $                true          USD 1410484909
## 6          $                true          USD 1414008752
##      disable_communication friends  fx_rate  goal      id is_backing
## 1          false            1.308394  1000 1361161119
## 2          false            1.000000  5000 746509287
## 3          false            1.000000  6000 1402909261
## 4          false            1.308394   400 311541751
## 5          false            1.000000 11000 466957735
## 6          false            1.000000  2000 1471254290
##      is_starrable is_starred launched_at      location
## 1          false            1517306009      London
## 2          false            1426520946      Columbus
## 3          false            1529070876 St. Petersburg
## 4          false            1519937886      Dorset
## 5          false            1407892909      Ypsilanti
## 6          false            1411416752 San Francisco
##                               name permissions pledged
## 1 Fine Bone China Ceramic Mugs, Made in England      1111.0
## 2           P.R.O.B.E. Magazine                      0.0
## 3           'Merican Wasteland Pilot                  8101.0
## 4           Cliff - Feature Film                     1116.5
## 5           Photo to Artwork                          11.0
## 6   Lilah Magazine 1st issue launching Dec 2014      826.0
##
## 1
## 2
## 3
## 4 {"id":3322408,"project_id":3322408,"state":"active","state_changed_at":1523464237,"name":"CLIFF - I
## 5

```

```
## 6
##
## 1 https://www.kickstarter.com/discover/categories/crafts
## 2 https://www.kickstarter.com/discover/categories/publishing/periodicals
## 3 https://www.kickstarter.com/discover/categories/film%20&%20video
## 4 https://www.kickstarter.com/discover/categories/film%20&%20video
## 5 https://www.kickstarter.com/discover/categories/art/digital%20art
## 6 https://www.kickstarter.com/discover/categories/publishing/periodicals
## spotlight staff_pick state state_changed_at static_usd_rate
## 1 true false successful 1521190409 1.413819
## 2 false false failed 1429112947 1.000000
## 3 true false successful 1531713540 1.000000
## 4 true false successful 1522443600 1.390235
## 5 false false failed 1410484909 1.000000
## 6 false false failed 1414008752 1.000000
## usd_pledged usd_type preparation_duration launch_duration
## 1 1570.753 international 1695248 3884400
## 2 0.000 international 158141 2592000
## 3 8101.000 international 3964815 2642664
## 4 1552.197 international 83846 2505714
## 5 11.000 domestic 546624 2592000
## 6 826.000 international 265954 2592000
## launch_duration_r created_at_readable deadline_readable
## 1 44 2018-01-10 13:59:21 2018-03-16 04:53:29
## 2 30 2015-03-14 15:53:25 2015-04-15 11:49:06
## 3 30 2018-04-30 12:34:21 2018-07-15 23:59:00
## 4 29 2018-02-28 16:40:40 2018-03-30 17:00:00
## 5 30 2014-08-06 13:31:25 2014-09-11 21:21:49
## 6 30 2014-09-19 14:19:58 2014-10-22 16:12:32
## launched_at_readable
## 1 2018-01-30 04:53:29
## 2 2015-03-16 11:49:06
## 3 2018-06-15 09:54:36
## 4 2018-03-01 15:58:06
## 5 2014-08-12 21:21:49
## 6 2014-09-22 16:12:32
```

#3. **Data exploration and Visualization


3.1 Summarise the number of projects for each status

```
status_projects <- clean_data %>%
  group_by(clean_data$state) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
head(status_projects)
```

```
## # A tibble: 5 x 2
##   `clean_data$state` count
##   <fct>             <int>
## 1 successful         4172
## 2 failed             2801
## 3 canceled           323
## 4 live              253
## 5 suspended          25
```

Plot the number of projects for each status

```
ggplot(status_projects, aes(reorder(status_projects$`clean_data$state`, -count), count, fill=count)) +
  ggtitle("Projects by Status") + xlab("Project Status") + ylab("Frequency") +
  geom_text(aes(label=count), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```




3.2 Summarise the number of projects for each category

```
catagory_projects <- clean_data %>%
  group_by(clean_data$category) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
head(catagory_projects)
```

```
## # A tibble: 6 x 2
##   `clean_data$category` count
##   <chr>                <int>
## 1 film & video          1102
## 2 music                 858
## 3 technology            795
## 4 art                   787
## 5 publishing            707
## 6 fashion               495
```

Plot the popularity of each category, which is determined by the number of projects

```
ggplot(catagory_projects, aes(reorder(catagory_projects$`clean_data$category`, -count), count, fill=count)) +
  ggtitle("Projects by Category") + xlab("Project Category") + ylab("Frequency") +
  geom_text(aes(label=count), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```




3.3 What types of projects are being funded?

```
pledged_category <- clean_data %>%
  group_by(clean_data$category) %>%
  summarise(total = sum(usd_pledged)) %>%
  arrange(desc(total))
head(pledged_category)
```

```
## # A tibble: 6 x 2
##   `clean_data$category`      total
##   <chr>                  <dbl>
## 1 technology             31373736.
## 2 film & video           10398557.
## 3 design                 8108525.
## 4 fashion                6797765.
## 5 games                  6465490.
## 6 music                  4988363.
```

Plot the amount pledged by each category

```
ggplot(pledged_category, aes(reorder(pledged_category`clean_data$category`, -total), total/1000000, fill=pledged_category`clean_data$category`)) +
  ggtitle("Total Amount Pledged by Category") + xlab("Project Category") +
  ylab("Amount Pledged (USD millions)") +
  geom_text(aes(label=paste0("$", round(total/1000000,1))), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```




3.4 How much is pledged per backer for each category?

```
pledged_avg_category <- clean_data %>%
  group_by(clean_data$category) %>%
  summarise(pledged = sum(usd_pledged), backers=sum(backers_count)) %>%
  mutate(avg = pledged/backers) %>%
  arrange(desc(avg))
head(pledged_avg_category)
```

```
## # A tibble: 6 x 4
##   `clean_data$category` pledged backers avg
##   <chr>                <dbl>    <int> <dbl>
## 1 technology          31373736.  224188 140.
## 2 design              8108525.   73647 110.
## 3 fashion             6797765.   63802 107.
## 4 film & video       10398557.  104058 99.9
## 5 food               2745202.   29880 91.9
## 6 photography        1620027.   18585 87.2
```

Plot the amount pledged per backer for each category


```
ggplot(pledged_avg_category, aes(reorder(pledged_avg_category$`clean_data$category`, -avg), avg, fill=avg)) +
  ggtitle("Average Amount Pledged per Backer") + xlab("Project Category") +
  ylab("Amount Pledged (USD)") +
  geom_text(aes(label=paste0("$", round(avg,2))), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```



3.5 Get the 10 highest goal successful projects

```
top_ten_success <- clean_data[clean_data$state == "successful",] %>%
  select("category", "goal", "state") %>%
  arrange(desc(goal))
head(top_ten_success)
```

```
##      category    goal    state
## 1  technology 1500000 successful
## 2  technology  800000 successful
## 3  technology  800000 successful
## 4 photography  700000 successful
## 5  technology  500000 successful
## 6      design  500000 successful
```

3.6 Get the average project goal


```
goal_avg <- clean_data %>%
  group_by(category) %>%
  summarise(goals = sum(goal), projects = n()) %>%
```

```
mutate(avg = goals/projects) %>%
  arrange(desc(avg))
head(goal_avg)
```

```
## # A tibble: 6 x 4
##   category      goals projects    avg
##   <chr>         <dbl>    <int>  <dbl>
## 1 design      104348985    417 250237.
## 2 film & video 150585489.   1102 136647.
## 3 technology   66183820    795  83250.
## 4 theater     9656489    221  43695.
## 5 journalism   6530680    170  38416.
## 6 food       16506111    458  36040.
```

Plot the average project goal.

```
ggplot(goal_avg, aes(reorder(goal_avg$category, -avg), avg, fill=avg)) + geom_bar(stat="identity") +
  ggtitle("Average Project Goal") + xlab("Project Category") + ylab("Project Goal (USD)") +
  geom_text(aes(label=paste0("$", round(avg,0))), vjust=-0.5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
        axis.text.x=element_text(size=12, angle=90), legend.position="null") +
  scale_fill_gradient(low="skyblue1", high="royalblue4")
```



3.7 percentage for projects in each category

```
perc_projects <- clean_data %>%
  filter(state %in% c("successful", "failed")) %>%
  group_by(category, state) %>%
```

```

summarize(count=n()) %>%
mutate(pct=count/sum(count)) %>%
arrange(desc(state), pct)
head(perc_projects)

```

```

## # A tibble: 6 x 4
## # Groups:   category [6]
##   category    state    count  pct
##   <chr>      <fct>    <int> <dbl>
## 1 journalism successful    43 0.283
## 2 food       successful   132 0.318
## 3 technology successful   344 0.464
## 4 crafts     successful   199 0.524
## 5 art        successful   390 0.536
## 6 design     successful   198 0.541

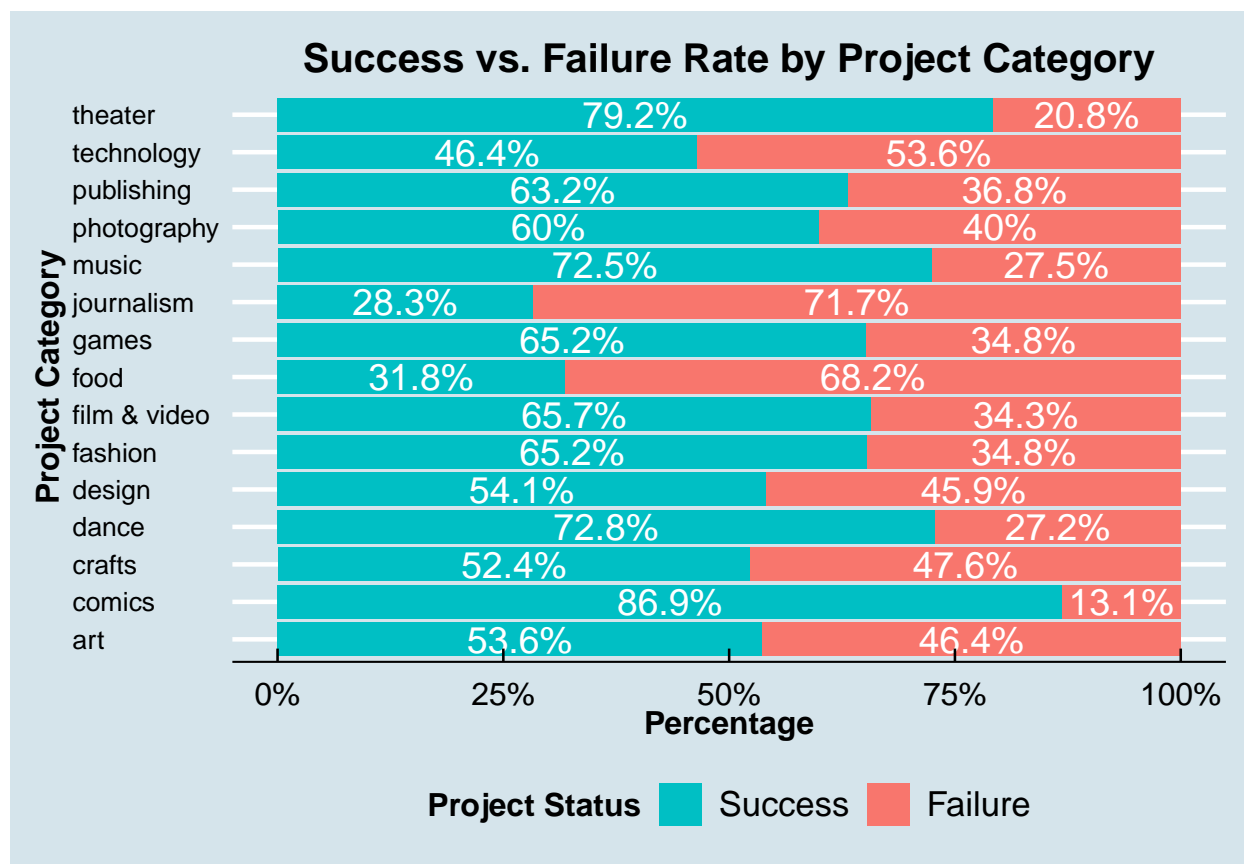
```

Plot the percentage for each category

```

ggplot(perc_projects, aes(perc_projects$category, pct, fill=state)) + geom_bar(stat="identity") +
  ggtitle("Success vs. Failure Rate by Project Category") +
  xlab("Project Category") + ylab("Percentage") + scale_y_continuous(labels=scales::percent) +
  scale_fill_discrete(name="Project Status", breaks=c("successful", "failed"),
    labels=c("Success", "Failure")) +
  geom_text(aes(label=paste0(round(pct*100,1),"%")), position=position_stack(vjust=0.5),
    colour="white", size=5) + theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"),
    axis.text.x=element_text(size=12), legend.position="bottom",
    legend.title=element_text(size=12, face="bold")) + coord_flip()

```




3.8 Does project length affect success rate?

```
perc_length <- clean_data %>%
  filter(state %in% c("successful", "failed"), launch_duration_r < 61) %>%
  group_by(launch_duration_r, state) %>%
  summarize(count=n()) %>%
  mutate(pct=count/sum(count))
head(perc_length)
```

```
## # A tibble: 6 x 4
## # Groups:   launch_duration_r [4]
##   launch_duration_r state      count    pct
##         <dbl> <fct>         <int> <dbl>
## 1             1 failed             1 0.5
## 2             1 successful           1 0.5
## 3             2 successful           1 1
## 4             3 failed             6 0.667
## 5             3 successful           3 0.333
## 6             4 failed             2 1
```

```
ggplot(perc_length[perc_length$state=="successful",], aes(launch_duration_r, pct)) +
  geom_point(colour="royalblue4", size=2.5) + ggtitle("Success Rate vs. Project Length") +
  xlab("Project Length (Days)") + ylab("Success Rate (%)") +
  scale_x_continuous(breaks=c(0,10,20,30,40,50,60)) + geom_vline(xintercept=30, colour="red") +
  theme_economist() +
  theme(plot.title=element_text(hjust=0.5), axis.title=element_text(size=12, face="bold"))
```



3.9 Where it is coming from?

```
countries_freq <- clean_data %>%
  group_by(country) %>%
  summarize(count=n())

countries.match <- joinCountryData2Map(countries_freq, joinCode="ISO2", nameJoinColumn="country")

## 22 codes from your data successfully matched countries in the map
## 0 codes from your data failed to match with a country code in the map
## 220 codes from the map weren't represented in your data

mapCountryData(countries.match, nameColumnToPlot="count",
  mapTitle="Number of Projects by Country", catMethod="logFixedWidth",
  colourPalette="heat")
```

Number of Projects by Country

